

both stages run on the OpenAI Agents SDK or the Claude Agent SDK — same prompts, same tools

Paper

PDF / MD / text

Stage 1 Harsh Critic

critical review of the paper

Stage 2 Merger

verify, anchor, calibrate

Final review

<score> <decision>

Stage 2 Merger

2a Filter

check every weakness
against the paper, drop
the ones that do not
hold, commit the draft

2b Anchor

retrieve reviewed papers
from every score band
and bracket the band
this paper belongs to

2c Refine

narrow inside that
bracket, then set the
score and decision
relative to the anchors

*calibration_search
up to 3 batched rounds*

Calibration corpus · 13k human-reviewed papers

vector retrieval, filtered to a band of real reviewer scores